

M1 — Replication and Interrogation

Group · Presented in class; deck and code zip due the same day · Weight and deadline: see the syllabus

What M1 is

Your five weekly project increments, compiled into **one replication-and-interrogation deck**, closing with your group’s judgment of the **single threat that most endangers the paper’s conclusion**.

Increment	Feeds M1 as
P1	Analytic cohort and eligibility
P2	Confounding, adjustment set, Table 1
P3	Effect modification
P4	Design-aware estimate
P5	Missing-data sensitivity

P1–P5 are ungraded on their own. M1 is where they are assessed.

What you submit

Two files, by the deadline in the syllabus:

1. **A slide deck (PDF)**. M1 is submitted strictly as a presentation deck. *Traditional written reports are not accepted for this milestone.*
2. **A zip archive of your reproducible repository, downloaded from GitHub**. You are **not** required to grant us access to the repository itself.

Your deck must include a **cover note** on the final slide, and the same text as `COVER-NOTE.md` in the root of your repository zip. It contains:

- an **AI Use Statement** (tool and version, where used, what you did with the output, and optionally representative prompts);
- a **one-line contribution note per member** — e.g. “*AB: cohort reconstruction, survey weighting code. CD: DAG, confounder set, Table 1. EF: slides, missing-data sensitivity.*”;
- any **external-feedback disclosure** — if a conversation with another group materially shaped an analytic decision, say so. That is a strength, not a confession.

The closing judgment

This is the part that distinguishes M1 from a summary of your increments.

Do not simply list the threats you found. **Rank one above the others**, and state the principle by which you ranked it — applied to the rival threats **by name**. A deck that lists threats without ranking them, or names a dominant threat without saying why it outranks the others, has not done this part of the task.

“The rest” means the threats **you** raised. You are not penalised for a threat you never mentioned under this heading — that shows up elsewhere.

Reproducibility

Every number, figure and table in your deck must have been produced by code in your repository, and that code must run start to finish on a fresh machine.

If you adapted the original authors' published code — which is legitimate replication work — **cite it in your repository README**, and identify in your cover note which scripts or functions are derived from it and what you changed.

How it is graded

M1 is graded on the coarse three-band scale used for M1–M4:

- **Exceptional** — deep insight, discrepancies traced to their root cause, or masterful engagement with counter-arguments.
- **Adequate** — the mechanics executed correctly, with sound epidemiological understanding.
- **Weak** — generic, superficial, or incorrect.

Hitting the paper's exact N is **not** the goal. A traced and explained discrepancy scores higher than an unexplained exact match.

What happens after you submit

Your deck and code zip are **released to your assigned reviewers**. Reviewer pairings are announced when M1 is submitted, and run in a cycle so that no group reviews the group reviewing it.

At **M2**, in the following class, **another group presents its critique of your M1 — not you**. You present your critique of the M1 *you* were assigned. Every group therefore does both: reviewed once, reviewer once.

Weight, deadline and policy

See the syllabus. It governs what this milestone is worth, when it is due, whether the weight can be shifted, and any late penalty. This page describes the task only.